

# R-CNN, Fast R-CNN, Faster R-CNN

Theoretical Part


17기 이형구

# Table of Contents

- 0. Introduction to Object Detection
- 1. R-CNN
- 2. SPP-Net
- 3. Fast R-CNN
- 4. Faster R-CNN

# Computer Vision task

- Types of Computer Vision
  - One Object: Classification, Localization
  - Multiple Objects: Detection, Segmentation



# Detection task

- Types of Detection
  - One-stage Detector
    - : Yolo, SSD, Retina-net
  - Two-stage detector
    - : R-CNN, SPP-Net, Fast R-CNN, Faster- RCNN, Pyramid Networks

# Object Detection


- Classification vs. Object detection

**Classification**



CAT


**Object Detection**



CAT, DOG, DUCK

# Prior Knowledge

- Selective Search
  - Bottom-up segmentation (merging regions at multiple scales)
- IOU (Intersect Over Union)




Intersection over union (IoU)

$$= \frac{\text{size of } \textcolor{yellow}{\square}}{\text{size of } \textcolor{blue}{\square}}$$

“Correct” if  $\text{IoU} \geq 0,5$

# 1. R-CNN




- Input -> Selective Search -> CNN -> Classifiers(SVM) -> Bounding-box regression

## R-CNN: *Regions with CNN features*



1. Input  
image



2. Extract region  
proposals (~2k)

3. Compute  
CNN features


4. Classify  
regions

R-CNN

# Characteristics of R-CNN

- Selective search
  - 2,000 region proposal
- Multi-stage
  - Convolutional fine tuning -> SVM classification -> Bounding-box regression
- Image -> Region -> Resize -> Convolution feature -> Classify


# Architecture of R-CNN



# Problems of R-CNN


- Slow train speed
  - Need to apply CNN for each 2000 region proposal
- CNN features are not updated in response to SVM and Bounding-box regressors
- Complex multi-stage training pipeline
  - Requires 84 hours using K40 GPU

## 2. SPP-Net (Spatial Pyramid Pooling)



- R-CNN
  - Image -> **Crop/Warp** -> Conv layers -> Fc layers -> Output
- SPP-Net
  - Image -> Conv layers -> **SPP** -> Fc layers -> Output
    - (Conv Features)
    - (Region)
    - (Classify)

# Architecture of SPP-Net




# Still takes a lot of time!


- Multi-stage training
- Both time complexity and space complexity are too high
- Object Detection is slow

### 3. Fast R-CNN

- Single-stage
- Higher performance
- Faster testing
- Training Step
  - Input – Image & Object proposal
  - Create Convolutional feature map using the input image
  - Create Feature Vector through ROI pooling layer
  - Detect object class & adjust bounding box





# Architecture of Fast R-CNN



# What is ROI pooling layer?

- Reduce vector size into a fixed size through max pooling only in the ROI region
- Each ROI ( $r, c, h, w$ ) -> ( $r, c$ ): top-left corner coordinate
- $5 \times 7 \rightarrow$  Max pooling ->  $3 \times 3$





## Multi-task loss

- 두 개의 output layer
- 분류 : 각 ROI별 Discrete probability distribution(전체 K+1 카테고리)
- 회귀 : bounding box regression  $t^k = (t_x^k, t_y^k, t_w^k, t_n^k)$
- u,v : ground truth class&target

$$L(p, u, t^u, v) = \underline{L_{\text{cls}}(p, u)} + \lambda[u \geq 1] \underline{L_{\text{loc}}(t^u, v)}$$

배경일 경우 0

$$L_{\text{cls}}(p, u) = -\log p_u$$

<분류>

$$L_{\text{loc}}(t^u, v) = \sum_{i \in \{\text{x}, \text{y}, \text{w}, \text{h}\}} \text{smooth}_{L_1}(t_i^u - v_i)$$
$$\text{smooth}_{L_1}(x) = \begin{cases} 0.5x^2 & \text{if } |x| < 1 \\ |x| - 0.5 & \text{otherwise} \end{cases}$$


<회귀>

# Mini-batch sampling

- N=2 mini batches
- R=128 (64 ROIs per image)
- 25% of ROI is positive sample ( $\text{IOU} \geq 0.5$ )
- $0.1 < \text{IOU} < 0.5$  indicates background (negative sample)


# Test Results

- Slow when testing
- Why? Region Proposal takes a long time



# Conclusion


- Faster and more accurate than R-CNN and SPP-Net
- One-stage structure
- But still not fast enough



RPN (Region Proposal Network)

# 4. Faster R-CNN


- Apply Region proposal using Neural Network (CPU -> GPU)
- Feature map can be used for generating region proposal




**RPN** + **Fast RCNN**

# Architecture of Faster R-CNN

- Insert a Region Proposal Network (RPN) after the last convolutional layer → using GPU!
- RPN trained to produce region proposals directly; no need for external region proposals
- After RPN, use RoI Pooling and an upstream classifier and bbox regressor just like Fast R-CNN



# Training Goal : Share Features



# Training Step

```
# Let M0 be an ImageNet pre-trained network

1. train_rpn(M0) → M1          # Train an RPN initialized from M0, get M1

2. generate_proposals(M1) → P1    # Generate training proposals P1 using RPN M1

3. train_fast_rcnn(M0, P1) → M2  # Train Fast R-CNN M2 on P1 initialized from M0

4. train_rpn_frozen_conv(M2) → M3 # Train RPN M3 from M2 without changing conv layers


5. generate_proposals(M3) → P2

6. train_fast_rcnn_frozen_conv(M3, P2) → M4 # Conv layers are shared with RPN M3


7. return add_rpn_layers(M4, M3.RPN)        # Add M3's RPN layers to Fast R-CNN M4
```

# RPN (Region Proposal Networks)

- Slide a small window on the feature map
- Build a small network for both
  - Regressing bounding-box locations
  - Classifying object




# RPN – Anchor



# RPN – Fully Convolutional Network

- Intermediate Layer – 256(or 512)  
3x3 filter, stride 1, padding 1
- Cls layer – 18(9x2) 1x1 filter, stride 1, padding 0
- Reg layer – 36(9x4) 1x1 filter, stride 1, padding 0



# RPN – Loss Function

$$L(\{p_i\}, \{t_i\}) = \frac{1}{N_{cls}} \sum_i L_{cls}(p_i, p_i^*) + \lambda \frac{1}{N_{reg}} \sum_i p_i^* L_{reg}(t_i, t_i^*)$$

Annotations for the equation:

- $i$  = anchor index in minibatch
- $L_{cls}(p_i, p_i^*)$ : Log loss
- $p_i^*$ : Ground truth objectness label
- $L_{reg}(t_i, t_i^*)$ : Smooth L1 loss
- $t_i^*$ : True box coordinates

Predicted probability of being an object for anchor  $i$

$N_{cls}$  = Number of anchors in minibatch (~ 256)

$N_{reg}$  = Number of anchor locations (~ 2400)

In practice  $\lambda = 10$ , so that both terms are roughly equally balanced

# Results

	R-CNN	Fast R-CNN	Faster R-CNN
Test time per image (with proposals)	50 seconds	2 seconds	<b>0.2 seconds</b>
(Speedup)	1x	25x	<b>250x</b>
mAP (VOC 2007)	66.0	<b>66.9</b>	<b>69.9</b>

Table 5: Timing (ms) on a K40 GPU, except SS proposal is evaluated in a CPU. “Region-wise” includes NMS, pooling, fully-connected, and softmax layers. See our released code for the profiling of running time.

model	system	conv	proposal	region-wise	total	rate
VGG	SS + Fast R-CNN	146	1510	174	1830	0.5 fps
VGG	RPN + Fast R-CNN	141	<b>10</b>	47	<b>198</b>	<b>5 fps</b>
ZF	RPN + Fast R-CNN	31	3	25	59	<b>17 fps</b>

# Experiments

Table 1: the learned average proposal size for each anchor using the ZF net (numbers for  $s = 600$ ).

anchor	$128^2, 2:1$	$128^2, 1:1$	$128^2, 1:2$	$256^2, 2:1$	$256^2, 1:1$	$256^2, 1:2$	$512^2, 2:1$	$512^2, 1:1$	$512^2, 1:2$
proposal	$188 \times 111$	$113 \times 114$	$70 \times 92$	$416 \times 229$	$261 \times 284$	$174 \times 332$	$768 \times 437$	$499 \times 501$	$355 \times 715$

Table 8: Detection results of Faster R-CNN on PASCAL VOC 2007 test set using **different settings of anchors**. The network is VGG-16. The training data is VOC 2007 trainval. The default setting of using 3 scales and 3 aspect ratios (69.9%) is the same as that in Table 3.

settings	anchor scales	aspect ratios	mAP (%)
1 scale, 1 ratio	$128^2$	1:1	65.8
	$256^2$	1:1	66.7
1 scale, 3 ratios	$128^2$	{2:1, 1:1, 1:2}	68.8
	$256^2$	{2:1, 1:1, 1:2}	67.9
3 scales, 1 ratio	{ $128^2, 256^2, 512^2$ }	1:1	69.8
3 scales, 3 ratios	{ $128^2, 256^2, 512^2$ }	{2:1, 1:1, 1:2}	69.9

Table 9: Detection results of Faster R-CNN on PASCAL VOC 2007 test set using **different values of  $\lambda$**  in Equation (1). The network is VGG-16. The training data is VOC 2007 trainval. The default setting of using  $\lambda = 10$  (69.9%) is the same as that in Table 3.

$\lambda$	0.1	1	10	100
mAP (%)	67.2	68.9	69.9	69.1

# Is it enough?

- ROI Pooling may introduce misalignments between ROI and extracted features
- Mask R-CNN

# References

- <https://youtu.be/kcPAGIgBGRs>
- <https://youtu.be/Jo32zrxr6l8>
- <https://youtu.be/HmJWvwlpW5g>